

# speech-recognition-assignment

May 15, 2023

Name1: Ali Hassan Ali Ahmed ElSharawy - ID1: 19016013

Name2: Ziad Mohamed Abuelkher - ID2: 19015733

Name3: Youssef Magdy Helmy - ID3: 19016937

```
[ ]: # This Python 3 environment comes with many helpful analytics libraries
      ↳ installed
      # It is defined by the kaggle/python Docker image: https://github.com/kaggle/
      ↳ docker-python
      # For example, here's several helpful packages to load

import numpy as np
import os
from scipy.io import wavfile
# ! pip install pydub
import matplotlib.pyplot as plt
import librosa.display
from pydub import AudioSegment
import IPython
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Conv1D, MaxPooling1D, Flatten,
      ↳ Dense, Dropout, Activation, BatchNormalization
from sklearn.preprocessing import StandardScaler
from tensorflow import keras
from tensorflow.keras.layers import *
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.regularizers import l2
```

walk through the dataset

```
[ ]: for dirname, _, filenames in os.walk('/kaggle/input'):
      for filename in filenames:
          print(os.path.join(dirname, filename))
      # You can write up to 20GB to the current directory (/kaggle/working/) that
      ↳ gets preserved as output when you create a version using "Save & Run All"
      # You can also write temporary files to /kaggle/temp/, but they won't be saved
      ↳ outside of the current session
```

useful function to treat with data

```
[3]: def read_audio(filepath):
      data, sampling_rate = librosa.load(filepath, sr=None)
      return data, sampling_rate

def plot_audio(filepath):
    data, sampling_rate = read_audio(filepath)
    plt.figure(figsize=(15, 5))
    plt.title('Waveform plot')
    plt.xlabel('Time (seconds)')
    plt.ylabel('Amplitude')
    librosa.display.waveshow(data, sr=sampling_rate)

def plot_audio_data(data):
    plt.figure(figsize=(15, 5))
    plt.title('Waveform plot')
    plt.xlabel('Time (seconds)')
    plt.ylabel('Amplitude')
    librosa.display.waveshow(data)

def play_audio(filepath):
    return IPython.display.Audio(filepath)
```

read files and sort directory list

```
[4]: path = '/kaggle/input/speech-emotion-recognition-en/Crema'
      dir_list = os.listdir(path)
      dir_list.sort()
      print(dir_list[0:10])
```

```
['1001_DFA_ANG_XX.wav', '1001_DFA_DIS_XX.wav', '1001_DFA_FEA_XX.wav',
'1001_DFA_HAP_XX.wav', '1001_DFA_NEU_XX.wav', '1001_DFA_SAD_XX.wav',
'1001_IEO_ANG_HI.wav', '1001_IEO_ANG_LO.wav', '1001_IEO_ANG_MD.wav',
'1001_IEO_DIS_HI.wav']
```

define emotions in data set

```
[5]: emotions = []
      for filename in dir_list:
          emotions.append(filename.split('_')[2])
      print(set(emotions))
```

```
{'NEU', 'DIS', 'ANG', 'FEA', 'HAP', 'SAD'}
```

we have 6 different emotions in our dataset: \* ANG : angry \* NEU : neutral \* HAP : happy \*  
SAD : sad \* DIS : disgust \* FEA : fear

map every emotion to number

```
[6]: dic = {'ANG' : 0, 'NEU' : 1, 'HAP' : 2, 'SAD' : 3, 'DIS' : 4, 'FEA' : 5}
```

read data from data set and store in lists : data, sample\_rate, labels and extract labels from files name


```
[7]: data = []
sample_rate = []
labels = []
for filename in dir_list:
    record, s = read_audio(path+'/'+filename)
    sample_rate.append(s)
    data.append(record)
    labels.append(dic[filename.split('_')[2]])
```

get familiar with data and try to change i to get different samples and look at different features

```
[8]: i = 1    # free to change it to display different audio and see wave form
p = path+'/'+dir_list[i]
emotion = dir_list[i].split('_')
print("emotion : ", emotion[2])
plot_audio(p)
IPython.display.Audio(p)
```

emotion : DIS

```
[8]: <IPython.lib.display.Audio object>
```




default rms in librosa package

```
[9]: energy = librosa.feature.rms(y=data[i])
plot_audio_data(energy)
print(energy.shape)
```

```
print(data[i].shape)
```


```
(1, 73)  
(37371,)
```



we need to get optimal values of frame\_length, hop\_length to get features of data good enough for training and don't lose information

```
[10]: energy = librosa.feature.rms(y=data[i],frame_length=2048,hop_length=200)  
plot_audio_data(energy)  
print(energy.shape)
```

```
(1, 187)
```

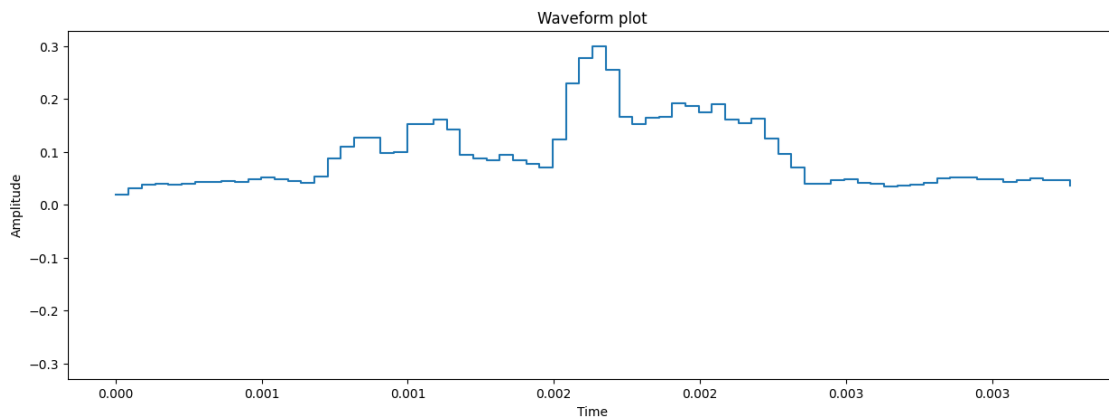


default zero crossing rate

```
[11]: zrc = librosa.feature.zero_crossing_rate(y=data[i])  
plot_audio_data(zrc)
```

```
print(zrc.shape)
```


(1, 73)



we need to get optimal values of frame\_length, hop\_length to get features of data good enough for training and don't lose information

```
[12]: zrc = librosa.feature.  
      ↪ zero_crossing_rate(y=data[i], frame_length=1000, hop_length=200)  
      plot_audio_data(zrc)  
      print(zrc.shape)
```

(1, 187)



melspectrogram in librosa package


```
[13]: S = librosa.feature.melspectrogram(y=data[i], n_mels=128, fmax=8000, ↪  
      ↪ n_fft=2048, hop_length=512)
```

```

# Convert to decibels
S_dB = librosa.power_to_db(S, ref=np.max)

# Display Mel spectrogram
librosa.display.specshow(S_dB, x_axis='time', y_axis='mel', fmax=8000)
plt.colorbar(format='%+2.0f dB')
plt.title('Mel-frequency spectrogram')
plt.show()
print(S_dB.shape)

```



(128, 73)

we need to make all data has same length to be able to train data and do different operation on it  
get max data file size to make all files has same length by padding we add zeros from back

```

[14]: max_len = 0
      min_len = float('inf')

      # loop over the list and update max and min lengths

```

```

for arr in data:
    length = len(arr)
    if length > max_len:
        max_len = length
    if length < min_len:
        min_len = length

print("Max length:", max_len)
print("Min length:", min_len)

```

Max length: 80080

Min length: 20287

```

[15]: padded_data = []
      for signal in data:
          pad_width = max_len - len(signal)
          padded_signal = np.pad(signal, (0, pad_width), mode='constant')
          padded_data.append(padded_signal)

```

we store data in padded form so all data has same size

```

[16]: IPython.display.Audio(padded_data[8], rate = 20000)

```

```

[16]: <IPython.lib.display.Audio object>

```

```

[17]: IPython.display.Audio(padded_data[13], rate = 20000)

```

```

[17]: <IPython.lib.display.Audio object>

```

Create Features Space from data

```

[18]: data_rms = []
      data_zcr = []
      data_mel = []

      for audio in padded_data:
          data_rms.append(librosa.feature.
              ↪rms(y=audio, frame_length=2048, hop_length=200))

      for audio in padded_data:
          data_zcr.append(librosa.feature.
              ↪zero_crossing_rate(y=audio, frame_length=1000, hop_length=200))

      for audio in padded_data:
          S = librosa.feature.melspectrogram(y=audio, n_mels=128, fmax=8000,
              ↪n_fft=2048, hop_length=512)
          S_db = librosa.power_to_db(S, ref=np.max)
          data_mel.append(S_db)

```

split data into training, test and validation

```
[19]: data_train_val_rms, data_test_rms, label_train_val_rms, label_test_rms =   
      ↪train_test_split(data_rms, labels, test_size=0.3, random_state=42)  
data_train_val_zcr, data_test_zcr, label_train_val_zcr, label_test_zcr =   
      ↪train_test_split(data_zcr, labels, test_size=0.3, random_state=42)  
data_train_val_mel, data_test_mel, label_train_val_mel, label_test_mel =   
      ↪train_test_split(data_mel, labels, test_size=0.3, random_state=42)  
  
data_train_rms, data_val_rms, label_train_rms, label_val_rms =   
      ↪train_test_split(data_train_val_rms, label_train_val_rms, test_size=0.05,   
      ↪random_state=42)  
data_train_zcr, data_val_zcr, label_train_zcr, label_val_zcr =   
      ↪train_test_split(data_train_val_zcr, label_train_val_zcr, test_size=0.05,   
      ↪random_state=42)  
data_train_mel, data_val_mel, label_train_mel, label_val_mel =   
      ↪train_test_split(data_train_val_mel, label_train_val_mel, test_size=0.05,   
      ↪random_state=42)
```

```
[20]: sample_train_val, sample_test, label_train_val_sample, label_test_sample =   
      ↪train_test_split(sample_rate, labels, test_size=0.3, random_state=42)  
sample_train, sample_val, label_train_sample, label_val_sample =   
      ↪train_test_split(sample_train_val, label_train_val_sample, test_size=0.05,   
      ↪random_state=42)
```

normalize data and convert to numpy array

```
[21]: data_train_rms = np.array(data_train_rms)  
data_test_rms = np.array(data_test_rms)  
label_train_rms = np.array(label_train_rms)  
label_test_rms = np.array(label_test_rms)  
data_val_rms = np.array(data_val_rms)  
label_val_rms = np.array(label_val_rms)  
  
data_train_zcr = np.array(data_train_zcr)  
data_test_zcr = np.array(data_test_zcr)  
label_train_zcr = np.array(label_train_zcr)  
label_test_zcr = np.array(label_test_zcr)  
data_val_zcr = np.array(data_val_zcr)  
label_val_zcr = np.array(label_val_zcr)  
  
train_data_mel = np.array(data_train_mel)  
data_test_mel = np.array(data_test_mel)  
train_label_mel = np.array(label_train_mel)  
label_test_mel = np.array(label_test_mel)  
cv_data_mel = np.array(data_val_mel)
```



```
cv_label_mel = np.array(label_val_mel)
```

we will use sample rate as feature

```
[22]: sample_train = np.array(sample_train)
      sample_val = np.array(sample_val)
      label_train_sample = np.array(label_train_sample)
      label_val_sample = np.array(label_val_sample)

      sample_test = np.array(sample_test)
      label_test_sample = np.array(label_test_sample)
```

```
[23]: print(sample_train.shape)
      print(sample_val.shape)
      print(sample_test.shape)
```

```
(4948,)
```

```
(261,)
```

```
(2233,)
```

normalize data

normalize rms feature

```
[24]: rms_train_reshape = data_train_rms.reshape(data_train_rms.
      ↪shape[0],data_train_rms.shape[2])
      rms_val_reshape = data_val_rms.reshape(data_val_rms.shape[0],data_val_rms.
      ↪shape[2])
      rms_test_reshape = data_test_rms.reshape(data_test_rms.shape[0],data_test_rms.
      ↪shape[2])

      scaler_rms = StandardScaler()
      data_train_rms_norm = scaler_rms.fit_transform(rms_train_reshape)
      data_train_rms_norm = scaler_rms.transform(rms_train_reshape)
      data_test_rms_norm = scaler_rms.transform(rms_test_reshape)
      data_val_rms_norm = scaler_rms.transform(rms_val_reshape)
```

```
[25]: print(data_train_rms_norm.shape)
      print(data_test_rms_norm.shape)
      print(data_val_rms_norm.shape)
```

```
(4948, 401)
```

```
(2233, 401)
```

```
(261, 401)
```

```
[26]: data_train_rms_norm = data_train_rms_norm.reshape((data_train_rms_norm.
      ↪shape[0],data_train_rms_norm.shape[1],1))
      data_test_rms_norm = data_test_rms_norm.reshape((data_test_rms_norm.
      ↪shape[0],data_test_rms_norm.shape[1],1))
```

```
data_val_rms_norm = data_val_rms_norm.reshape((data_val_rms_norm.  
↪shape[0],data_val_rms_norm.shape[1],1))
```

normalize zcr feature

```
[27]: zcr_train_reshape = data_train_zcr.reshape(data_train_zcr.  
↪shape[0],data_train_zcr.shape[2])  
zcr_val_reshape = data_val_zcr.reshape(data_val_zcr.shape[0],data_val_zcr.  
↪shape[2])  
zcr_test_reshape = data_test_zcr.reshape(data_test_zcr.shape[0],data_test_zcr.  
↪shape[2])  
  
scaler_zcr = StandardScaler()  
data_train_zcr_norm = scaler_zcr.fit_transform(zcr_train_reshape)  
data_train_zcr_norm = scaler_zcr.transform(zcr_train_reshape)  
data_test_zcr_norm = scaler_zcr.transform(zcr_test_reshape)  
data_val_zcr_norm = scaler_zcr.transform(zcr_val_reshape)
```

```
[28]: print(data_train_zcr_norm.shape)  
print(data_test_zcr_norm.shape)  
print(data_val_zcr_norm.shape)
```

```
(4948, 401)  
(2233, 401)  
(261, 401)
```

```
[29]: data_train_zcr_norm = data_train_zcr_norm.reshape((data_train_zcr_norm.  
↪shape[0],data_train_zcr_norm.shape[1],1))  
data_test_zcr_norm = data_test_zcr_norm.reshape((data_test_zcr_norm.  
↪shape[0],data_test_zcr_norm.shape[1],1))  
data_val_zcr_norm = data_val_zcr_norm.reshape((data_val_zcr_norm.  
↪shape[0],data_val_zcr_norm.shape[1],1))
```

```
[30]: print(data_train_zcr_norm.shape)  
print(data_test_zcr_norm.shape)  
print(data_val_zcr_norm.shape)
```

```
(4948, 401, 1)  
(2233, 401, 1)  
(261, 401, 1)
```

normalize mel spectrogram feature

```
[33]: mel_train_reshape = train_data_mel.reshape(train_data_mel.  
↪shape[0],train_data_mel.shape[1]*train_data_mel.shape[2])  
mel_cv_reshape = cv_data_mel.reshape(cv_data_mel.shape[0],cv_data_mel.  
↪shape[1]*cv_data_mel.shape[2])
```

```
mel_test_reshape = data_test_mel.reshape(data_test_mel.shape[0],data_test_mel.
↪shape[1]*data_test_mel.shape[2])
print(cv_data_mel.shape)
```

(261, 128, 157)

```
[34]: scaler_mel = StandardScaler()
data_train_mel_norm = scaler_mel.fit_transform(mel_train_reshape)
data_train_mel_norm = scaler_mel.transform(mel_train_reshape)
data_test_mel_norm = scaler_mel.transform(mel_test_reshape)
data_cv_mel_norm = scaler_mel.transform(mel_cv_reshape)
```

```
[35]: print(data_train_mel_norm.shape)
print(data_test_mel_norm.shape)
print(data_cv_mel_norm.shape)
data_train_mel_norm = data_train_mel_norm.reshape((data_train_mel_norm.
↪shape[0],train_data_mel.shape[1],train_data_mel.shape[2]))
data_test_mel_norm = data_test_mel_norm.reshape((data_test_mel_norm.
↪shape[0],data_test_mel.shape[1], data_test_mel.shape[2]))
data_cv_mel_norm = data_cv_mel_norm.reshape((data_cv_mel_norm.
↪shape[0],cv_data_mel.shape[1],cv_data_mel.shape[2]))
print(data_train_mel_norm.shape)
print(data_test_mel_norm.shape)
print(data_cv_mel_norm.shape)
```

(4948, 20096)

(2233, 20096)

(261, 20096)

(4948, 128, 157)

(2233, 128, 157)

(261, 128, 157)

concatenate data to get time domain feature

```
[36]: train_time = np.concatenate((data_train_rms_norm, data_train_zcr_norm), axis = -1
↪)
val_time = np.concatenate((data_val_rms_norm, data_val_zcr_norm), axis = -1)
test_time = np.concatenate((data_test_rms_norm, data_test_zcr_norm), axis = -1)
```

```
[37]: train_time = train_time.reshape(train_time.shape[0], -1, 1)
val_time = val_time.reshape(val_time.shape[0], -1, 1)
test_time = test_time.reshape(test_time.shape[0], -1, 1)
```

add sample rate as second channel

```
[ ]: # sample_train = sample_train.reshape(-1,1,1)
# sample_test = sample_test.reshape(-1,1,1)
# sample_val = sample_val.reshape(-1,1,1)
```

```
[ ]: # print(sample_train.shape)
      # print(train_time.shape)
```

```
[ ]: # train_time = np.concatenate((train_time, sample_train), axis = 1)
      # val_time = np.concatenate((val_time, sample_val), axis = 1)
      # test_time = np.concatenate((test_time, sample_test), axis = 1)
```

```
[39]: # print(train_time.shape)
       # print(label_time_train.shape)
```

RMSE model

build model

```
[40]: model_time = Sequential()
      alpha = 0.000001
      model_time.add(Conv1D(64, 16, padding='same', input_shape=(train_time.
        ↪shape[1],1)))

      model_time.add(LeakyReLU(alpha=alpha))
      model_time.add(BatchNormalization())

      model_time.add(Conv1D(64, 16, padding='same'))

      model_time.add(LeakyReLU(alpha=alpha))
      model_time.add(Conv1D(64, 16, padding='same'))

      model_time.add(LeakyReLU(alpha=alpha))

      model_time.add(Conv1D(64, 16, padding='same'))

      model_time.add(LeakyReLU(alpha=alpha))
      model_time.add(Conv1D(64, 16, padding='same'))

      model_time.add(LeakyReLU(alpha=alpha))

      model_time.add(MaxPooling1D(pool_size=(8)))
      model_time.add(Conv1D(64, 16, padding='same'))
      model_time.add(LeakyReLU(alpha=alpha))

      model_time.add(Conv1D(64, 8, padding='same'))
      model_time.add(LeakyReLU(alpha=alpha))

      model_time.add(BatchNormalization())
      model_time.add(Conv1D(128, 8, padding='same'))
      model_time.add(LeakyReLU(alpha=alpha))
      model_time.add(Conv1D(128, 8, padding='same'))
      model_time.add(LeakyReLU(alpha=alpha))
```

```

model_time.add(Dropout(0.2))
model_time.add(Conv1D(256, 8, padding='same'))
model_time.add(LeakyReLU(alpha=alpha))

model_time.add(MaxPooling1D(pool_size=(8)))
model_time.add(BatchNormalization())
model_time.add(Conv1D(256, 8, padding='same'))
model_time.add(LeakyReLU(alpha=alpha))

model_time.add(Conv1D(256, 4, padding='same'))
model_time.add(LeakyReLU(alpha=alpha))

model_time.add(Conv1D(512, 4, padding='same'))
model_time.add(LeakyReLU(alpha=alpha))
model_time.add(Conv1D(512, 4, padding='same'))
model_time.add(LeakyReLU(alpha=alpha))

model_time.add(Flatten())
model_time.add(Dense(128, activation='relu', kernel_regularizer=l2(0.01)))
model_time.add(Dense(64, activation='relu', kernel_regularizer=l2(0.01)))
model_time.add(Dense(6, activation='softmax'))
model_time.summary()

```

Model: "sequential"

| Layer (type)                              | Output Shape    | Param # |
|---|-----------------|---------|
| conv1d (Conv1D)                           | (None, 802, 64) | 1088    |
| leaky_re_lu (LeakyReLU)                   | (None, 802, 64) | 0       |
| batch_normalization (Batch Normalization) | (None, 802, 64) | 256     |
| conv1d_1 (Conv1D)                         | (None, 802, 64) | 65600   |
| leaky_re_lu_1 (LeakyReLU)                 | (None, 802, 64) | 0       |
| conv1d_2 (Conv1D)                         | (None, 802, 64) | 65600   |
| leaky_re_lu_2 (LeakyReLU)                 | (None, 802, 64) | 0       |
| conv1d_3 (Conv1D)                         | (None, 802, 64) | 65600   |
| leaky_re_lu_3 (LeakyReLU)                 | (None, 802, 64) | 0       |

|   |                  |        |
|---|------------------|--------|
| conv1d_4 (Conv1D)                           | (None, 802, 64)  | 65600  |
| leaky_re_lu_4 (LeakyReLU)                   | (None, 802, 64)  | 0      |
| max_pooling1d (MaxPooling1D)                | (None, 100, 64)  | 0      |
| conv1d_5 (Conv1D)                           | (None, 100, 64)  | 65600  |
| leaky_re_lu_5 (LeakyReLU)                   | (None, 100, 64)  | 0      |
| conv1d_6 (Conv1D)                           | (None, 100, 64)  | 32832  |
| leaky_re_lu_6 (LeakyReLU)                   | (None, 100, 64)  | 0      |
| batch_normalization_1 (Batch Normalization) | (None, 100, 64)  | 256    |
| conv1d_7 (Conv1D)                           | (None, 100, 128) | 65664  |
| leaky_re_lu_7 (LeakyReLU)                   | (None, 100, 128) | 0      |
| conv1d_8 (Conv1D)                           | (None, 100, 128) | 131200 |
| leaky_re_lu_8 (LeakyReLU)                   | (None, 100, 128) | 0      |
| dropout (Dropout)                           | (None, 100, 128) | 0      |
| conv1d_9 (Conv1D)                           | (None, 100, 256) | 262400 |
| leaky_re_lu_9 (LeakyReLU)                   | (None, 100, 256) | 0      |
| max_pooling1d_1 (MaxPooling1D)              | (None, 12, 256)  | 0      |
| batch_normalization_2 (Batch Normalization) | (None, 12, 256)  | 1024   |
| conv1d_10 (Conv1D)                          | (None, 12, 256)  | 524544 |
| leaky_re_lu_10 (LeakyReLU)                  | (None, 12, 256)  | 0      |
| conv1d_11 (Conv1D)                          | (None, 12, 256)  | 262400 |
| leaky_re_lu_11 (LeakyReLU)                  | (None, 12, 256)  | 0      |
| conv1d_12 (Conv1D)                          | (None, 12, 512)  | 524800 |

|                            |                 |         |
|----------------------------|-----------------|---------|
| leaky_re_lu_12 (LeakyReLU) | (None, 12, 512) | 0       |
| conv1d_13 (Conv1D)         | (None, 12, 512) | 1049088 |
| leaky_re_lu_13 (LeakyReLU) | (None, 12, 512) | 0       |
| flatten (Flatten)          | (None, 6144)    | 0       |
| dense (Dense)              | (None, 128)     | 786560  |
| dense_1 (Dense)            | (None, 64)      | 8256    |
| dense_2 (Dense)            | (None, 6)       | 390     |

```
=====
Total params: 3,978,758
Trainable params: 3,977,990
Non-trainable params: 768
-----
```

concatenate label data

```
[41]: from keras.utils import to_categorical

label_time_train = to_categorical(label_train_rms, 6)
label_time_test = to_categorical(label_test_rms, 6)
label_time_val = to_categorical(label_val_rms, 6)
```

```
[42]: print(train_time.shape)
print(val_time.shape)
```

```
(4948, 802, 1)
(261, 802, 1)
```

```
[44]: from tensorflow.keras.callbacks import EarlyStopping

early_stopping_callback = EarlyStopping(monitor='val_accuracy', patience=5)

model_time.compile(loss='categorical_crossentropy', optimizer='adam',
    ↪metrics=['accuracy'])
model_history=model_time.fit(train_time, label_time_train, batch_size=32,
    ↪epochs=100,
                                validation_data=(val_time, label_time_val),
    ↪callbacks=[early_stopping_callback])

# model.compile(loss='categorical_crossentropy', optimizer='adam',
    ↪metrics=['accuracy'])
```

```
# model_history=model.fit(train_time, label_time_train, batch_size=32,
↪epochs=20,
#                               validation_data=(val_time, label_time_val))
```

Epoch 1/100

155/155 [=====] - 18s 31ms/step - loss: 2.3396 -  
accuracy: 0.3333 - val\_loss: 1.9958 - val\_accuracy: 0.2605

Epoch 2/100

155/155 [=====] - 4s 23ms/step - loss: 1.6364 -  
accuracy: 0.3824 - val\_loss: 1.6556 - val\_accuracy: 0.3678

Epoch 3/100

155/155 [=====] - 4s 24ms/step - loss: 1.5467 -  
accuracy: 0.4131 - val\_loss: 1.5220 - val\_accuracy: 0.3908

Epoch 4/100

155/155 [=====] - 4s 24ms/step - loss: 1.4974 -  
accuracy: 0.4258 - val\_loss: 1.4430 - val\_accuracy: 0.4444

Epoch 5/100

155/155 [=====] - 4s 24ms/step - loss: 1.4798 -  
accuracy: 0.4274 - val\_loss: 1.4704 - val\_accuracy: 0.4521

Epoch 6/100

155/155 [=====] - 4s 24ms/step - loss: 1.4560 -  
accuracy: 0.4323 - val\_loss: 1.4050 - val\_accuracy: 0.4598

Epoch 7/100

155/155 [=====] - 4s 24ms/step - loss: 1.4322 -  
accuracy: 0.4351 - val\_loss: 1.5785 - val\_accuracy: 0.4406

Epoch 8/100

155/155 [=====] - 4s 24ms/step - loss: 1.4189 -  
accuracy: 0.4434 - val\_loss: 1.4125 - val\_accuracy: 0.4751

Epoch 9/100

155/155 [=====] - 4s 24ms/step - loss: 1.4139 -  
accuracy: 0.4378 - val\_loss: 1.3525 - val\_accuracy: 0.4368

Epoch 10/100

155/155 [=====] - 4s 24ms/step - loss: 1.4044 -  
accuracy: 0.4460 - val\_loss: 1.4105 - val\_accuracy: 0.4291

Epoch 11/100

155/155 [=====] - 4s 23ms/step - loss: 1.3930 -  
accuracy: 0.4483 - val\_loss: 1.3651 - val\_accuracy: 0.4674

Epoch 12/100

155/155 [=====] - 4s 23ms/step - loss: 1.3878 -  
accuracy: 0.4507 - val\_loss: 1.5386 - val\_accuracy: 0.4291


Epoch 13/100

155/155 [=====] - 4s 24ms/step - loss: 1.3822 -  
accuracy: 0.4535 - val\_loss: 1.3412 - val\_accuracy: 0.4559

evalute model performance



```
[45]: plt.plot(model_history.history['loss'])
plt.plot(model_history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
[47]: pred = model_time.predict(test_time)
```

70/70 [=====] - 1s 9ms/step

```
[48]: print(pred)
```

```
[[6.54770017e-01 6.02504704e-04 1.72963589e-01 4.95522516e-04
 2.90141385e-02 1.42154217e-01]
 [1.96295306e-02 2.63427854e-01 1.15780577e-01 2.98779100e-01
 1.39243722e-01 1.63139209e-01]
 [3.87055129e-02 1.47476450e-01 1.86895311e-01 1.10940345e-01
 2.53588766e-01 2.62393653e-01]]
```

```
...
[1.71514899e-02 3.19855750e-01 1.08466461e-01 2.87581831e-01
 1.34053141e-01 1.32891387e-01]
[2.88596991e-02 3.50182027e-01 1.53114676e-01 1.41299054e-01
 1.36858702e-01 1.89685851e-01]
[7.56240590e-03 5.93177676e-02 5.88369817e-02 3.32595825e-01
 2.77759373e-01 2.63927639e-01]]
```

```
[50]: pred_time = np.argmax(pred, axis=1)

print(pred_time)
label_time_test_time = np.argmax(label_time_test, axis=1)
print(label_time_test_time)
```

```
[0 3 5 ... 1 1 3]
[5 3 4 ... 3 1 4]
```

```
[51]: from sklearn.metrics import f1_score

f1 = f1_score(label_time_test_time, pred_time, average='weighted')
print("f score measure : ", f1)
```

```
f score measure : 0.4329533463971103
```

```
[73]: print(label_time_test)
```

```
[[0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 1. 0. 0.]
 [0. 0. 0. 0. 1. 0.]
 ...
 [0. 0. 0. 1. 0. 0.]
 [0. 1. 0. 0. 0. 0.]
 [0. 0. 0. 0. 1. 0.]]
```

```
[53]: score = model_time.evaluate(test_time, label_time_test, verbose=0)
print("accuracy measure : ",score[1])
print("loss measure : ",score[0])
```

```
accuracy measure : 0.45991939306259155
loss measure : 1.3897926807403564
```

```
=====
```

```
mel spectrogram mdl
```

```
[59]: model_mel = Sequential()
model_mel.add(Conv2D(32, kernel_size=(3, 3), activation='relu',
    ↪input_shape=(data_train_mel_norm.shape[1],data_train_mel_norm.shape[2],1)))
model_mel.add(MaxPooling2D(pool_size=(2, 2)))
```

```

model_mel.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
model_mel.add(MaxPooling2D(pool_size=(2, 2)))
model_mel.add(Conv2D(128, kernel_size=(3, 3), activation='relu'))
model_mel.add(MaxPooling2D(pool_size=(2, 2)))
model_mel.add(Flatten())
model_mel.add(Dense(128, activation='relu'))
model_mel.add(Dense(64, activation='relu'))
model_mel.add(Dense(6, activation='softmax'))
model_mel.summary()

```

Model: "sequential\_3"

| Layer (type)                    | Output Shape         | Param # |
|---------------------------------|----------------------|---------|
| conv2d_6 (Conv2D)               | (None, 126, 155, 32) | 320     |
| max_pooling2d_6 (MaxPooling 2D) | (None, 63, 77, 32)   | 0       |
| conv2d_7 (Conv2D)               | (None, 61, 75, 64)   | 18496   |
| max_pooling2d_7 (MaxPooling 2D) | (None, 30, 37, 64)   | 0       |
| conv2d_8 (Conv2D)               | (None, 28, 35, 128)  | 73856   |
| max_pooling2d_8 (MaxPooling 2D) | (None, 14, 17, 128)  | 0       |
| flatten_3 (Flatten)             | (None, 30464)        | 0       |
| dense_9 (Dense)                 | (None, 128)          | 3899520 |
| dense_10 (Dense)                | (None, 64)           | 8256    |
| dense_11 (Dense)                | (None, 6)            | 390     |
| Total params: 4,000,838         |                      |         |
| Trainable params: 4,000,838     |                      |         |
| Non-trainable params: 0         |                      |         |

convert to hot code label

```

[60]: from keras.utils import to_categorical
cat_train_label = to_categorical(train_label_mel, 6)
cat_cv_label = to_categorical(cv_label_mel)

```

```
print(cat_train_label.shape)
print(cat_cv_label.shape)
```

(4948, 6)

(261, 6)

train model

```
[61]: from tensorflow.keras.callbacks import EarlyStopping
early_stopping_callback = EarlyStopping(monitor='val_accuracy', patience=10)
model_mel.compile(loss='categorical_crossentropy', optimizer='adam',
    ↳metrics=['accuracy'])
model_history=model_mel.fit(data_train_mel_norm, cat_train_label,
    ↳batch_size=32, epochs=100, validation_data=(data_cv_mel_norm,
    ↳cat_cv_label),callbacks=[early_stopping_callback])
```

Epoch 1/100

155/155 [=====] - 4s 15ms/step - loss: 1.5282 -  
accuracy: 0.3771 - val\_loss: 1.4126 - val\_accuracy: 0.4789

Epoch 2/100

155/155 [=====] - 2s 15ms/step - loss: 1.3358 -  
accuracy: 0.4733 - val\_loss: 1.3421 - val\_accuracy: 0.5057

Epoch 3/100

155/155 [=====] - 2s 13ms/step - loss: 1.1980 -  
accuracy: 0.5325 - val\_loss: 1.1746 - val\_accuracy: 0.5517

Epoch 4/100

155/155 [=====] - 2s 13ms/step - loss: 1.0401 -  
accuracy: 0.6051 - val\_loss: 1.2360 - val\_accuracy: 0.5249

Epoch 5/100

155/155 [=====] - 2s 13ms/step - loss: 0.8521 -  
accuracy: 0.6750 - val\_loss: 1.1646 - val\_accuracy: 0.5939

Epoch 6/100

155/155 [=====] - 2s 13ms/step - loss: 0.6420 -  
accuracy: 0.7565 - val\_loss: 1.4382 - val\_accuracy: 0.5096

Epoch 7/100

155/155 [=====] - 2s 13ms/step - loss: 0.4157 -  
accuracy: 0.8486 - val\_loss: 1.5988 - val\_accuracy: 0.5441

Epoch 8/100

155/155 [=====] - 2s 13ms/step - loss: 0.2400 -  
accuracy: 0.9129 - val\_loss: 1.8979 - val\_accuracy: 0.5517

Epoch 9/100

155/155 [=====] - 2s 13ms/step - loss: 0.1266 -  
accuracy: 0.9572 - val\_loss: 2.1149 - val\_accuracy: 0.5134

Epoch 10/100


155/155 [=====] - 2s 13ms/step - loss: 0.1017 -  
accuracy: 0.9658 - val\_loss: 2.3297 - val\_accuracy: 0.5517

Epoch 11/100

155/155 [=====] - 2s 13ms/step - loss: 0.0871 -

```
accuracy: 0.9701 - val_loss: 2.6384 - val_accuracy: 0.5364
Epoch 12/100
155/155 [=====] - 2s 13ms/step - loss: 0.0824 -
accuracy: 0.9733 - val_loss: 2.7644 - val_accuracy: 0.5249
Epoch 13/100
155/155 [=====] - 2s 13ms/step - loss: 0.0580 -
accuracy: 0.9826 - val_loss: 2.9949 - val_accuracy: 0.5211
Epoch 14/100
155/155 [=====] - 2s 15ms/step - loss: 0.0475 -
accuracy: 0.9840 - val_loss: 2.9795 - val_accuracy: 0.5709
Epoch 15/100
155/155 [=====] - 2s 14ms/step - loss: 0.0538 -
accuracy: 0.9844 - val_loss: 3.0624 - val_accuracy: 0.5211
```

```
[62]: plt.plot(model_history.history['loss'])
plt.plot(model_history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



evalute model performance

```
[64]: pred = model_mel.predict(data_test_mel_norm)
      print(pred)
```

```
70/70 [=====] - 0s 6ms/step
[[9.99955535e-01 1.13819287e-22 3.78458826e-05 1.54699739e-13
  2.75636565e-16 6.62365983e-06]
 [1.80216812e-04 5.32235026e-01 8.40027269e-07 4.66214150e-01
  2.08244001e-05 1.34900946e-03]
 [7.17388320e-05 4.05109931e-05 5.35217300e-03 4.60997717e-06
  1.20479278e-01 8.74051690e-01]
 ...
 [2.83548841e-04 3.98541195e-03 8.54643166e-01 1.40758500e-01
  2.91849894e-04 3.74945193e-05]
 [1.21240455e-05 9.94802952e-01 1.22962518e-07 5.18436916e-03
  3.95581793e-07 8.43647499e-08]
 [2.14086685e-04 7.08788320e-06 3.71561535e-02 1.01014669e-03
  8.66688728e-01 9.49238539e-02]]
```

```
[67]: pred_mel = np.argmax(pred, axis=1)

      print(pred_mel)
      print(label_test_mel)
```

```
[0 1 5 ... 2 1 4]
[5 3 4 ... 3 1 4]
```

```
[69]: f1 = f1_score(label_test_mel, pred_mel, average='weighted')
      print("f score measure : ", f1)
```

```
f score measure : 0.5342704134448267
```

```
[74]: print(to_categorical(label_test_mel, 6))
```

```
[[0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 1. 0. 0.]
 [0. 0. 0. 0. 1. 0.]
 ...
 [0. 0. 0. 1. 0. 0.]
 [0. 1. 0. 0. 0. 0.]
 [0. 0. 0. 0. 1. 0.]]
```

```
[75]: score = model_mel.evaluate(data_test_mel_norm, to_categorical(label_test_mel),
      ↪ verbose=0)
      print("accuracy measure : ",score[1])
      print("loss measure : ",score[0])
```

```
accuracy measure : 0.5396327972412109
loss measure : 3.205275058746338
```

=====

confusion matrix

```
[76]: from sklearn.metrics import confusion_matrix

def get_confusion_matrix(actual, predicted):
    # create confusion matrix using sklearn's confusion_matrix function
    cm = confusion_matrix(actual, predicted)

    print(cm)
    return cm

def confusing_classes(conf_matrix, threshold, categories):
    pairs_numbers = []

    for i in range(len(conf_matrix)):
        for j in range(len(conf_matrix[i])):
            if i != j:
                misclassifications = conf_matrix[i][j] + conf_matrix[j][i]

                if misclassifications > threshold:
                    pairs_numbers.append(((categories[i], categories[j]),
↪misclassifications))

    pairs_numbers.sort(key=lambda x: x[1], reverse=True)

    return pairs_numbers

def plot_confusion_matrix(cm, classes):
    plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
    plt.title('Confusion matrix')
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes)
    plt.yticks(tick_marks, classes)
    fmt = '.2f'
    thresh = cm.max() / 2.
    for i, j in np.ndindex(cm.shape):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.tight_layout()
```

```
plt.show()
```

confusion matrix for time feature


```
[77]: conf_matrix = get_confusion_matrix(label_time_test_time, pred_time)
      classes = [0, 1, 2, 3, 4, 5]

      reverse_dic = {value: key for key, value in dic.items()}
      category_list = [reverse_dic[value] for value in classes]
      # print(category_list)

      plot_confusion_matrix(conf_matrix, category_list)
      threshold=100
      confused_classes = confusing_classes(conf_matrix, threshold, category_list)
      print(f"Confused classes with more than {threshold} misclassifications are:")
      i=0
      for pair, value in confused_classes:
          if(i%2==0):
              print(f"Classes ({pair[0]} , {pair[1]}) --> {value}")
          i+=1
```

```
[[291   6  41   3   9   9]
 [  7 144  62  65  33  42]
 [ 97  45 150  24  22  40]
 [  2  31  12 286  23  32]
 [ 54  32  43 123  72  62]
 [ 60  31  61 108  27  84]]
```





Confused classes with more than 100 misclassifications are:

Classes (SAD , DIS) --> 146

Classes (SAD , FEA) --> 140

Classes (ANG , HAP) --> 138

Classes (NEU , HAP) --> 107

Classes (HAP , FEA) --> 101

confusion matrix for mel spectrogram

```
conf_matrix = get_confusion_matrix(label_test_mel, pred_mel) classes = [0, 1, 2, 3, 4, 5]
```

```
reverse_dic = {value: key for key, value in dic.items()} category_list = [reverse_dic[value] for  
value in classes] # print(category_list)
```

```
plot_confusion_matrix(conf_matrix, category_list) threshold=100 confused_classes = confus-  
ing_classes(conf_matrix, threshold, category_list) print(f"Confused classes with more than  
{threshold} misclassifications are:") i=0 for pair, value in confused_classes: if(i%2==0):  
print(f"Classes ({pair[0]} , {pair[1]}) -> {value}") i+=1
```